

# Homework3

August 3, 2022


Homework 3

Compiled with

\$ jupyter nbconvert --to pdf Homework3.

## 1 Neural Networks

In this problem we will analyze a simple neural network to understand its classification properties. Consider the neural network given in the figure below, with ReLU activation functions (denoted by  $f$ ) on all neurons, and a softmax activation function in the output layer:



Given an input  $x = [x_1, x_2]^T$ , the hidden units in the network are activated in stages as described by the following equations:

$$\begin{aligned}
z_1 &= x_1 W_{11} + x_2 W_{21} + W_{01} & f(z_1) &= \max\{z_1, 0\} \\
z_2 &= x_1 W_{12} + x_2 W_{22} + W_{02} & f(z_2) &= \max\{z_2, 0\} \\
z_3 &= x_1 W_{13} + x_2 W_{23} + W_{03} & f(z_3) &= \max\{z_3, 0\} \\
z_4 &= x_1 W_{14} + x_2 W_{24} + W_{04} & f(z_4) &= \max\{z_4, 0\}
\end{aligned}$$

$$\begin{aligned}
u_1 &= f(z_1)V_{11} + f(z_2)V_{21} + f(z_3)V_{31} + f(z_4)V_{41} + V_{01} & f(u_1) &= \max\{u_1, 0\} \\
u_2 &= f(z_1)V_{12} + f(z_2)V_{22} + f(z_3)V_{32} + f(z_4)V_{42} + V_{02} & f(u_2) &= \max\{u_2, 0\}
\end{aligned}$$

The final output of the network is obtained by applying the softmax function to the last hidden layer,

$$\begin{aligned}
o_1 &= \frac{e^{f(u_1)}}{e^{f(u_1)} + e^{f(u_2)}} \\
o_2 &= \frac{e^{f(u_2)}}{e^{f(u_1)} + e^{f(u_2)}}
\end{aligned}$$

In this problem we will consider the following set of parameters:

$$\begin{bmatrix} W_{11} & W_{21} & W_{01} \\ W_{12} & W_{22} & W_{02} \\ W_{13} & W_{23} & W_{03} \\ W_{14} & W_{24} & W_{04} \end{bmatrix} = \begin{bmatrix} 1 & 0 & -1 \\ 0 & 1 & -1 \\ -1 & 0 & -1 \\ 0 & -1 & -1 \end{bmatrix}$$

$$\begin{bmatrix} V_{11} & V_{21} & V_{31} & V_{41} & V_{01} \\ V_{12} & V_{22} & V_{32} & V_{42} & V_{02} \end{bmatrix} = \begin{bmatrix} 1 & 1 & 1 & 1 & 0 \\ -1 & -1 & -1 & -1 & 2 \end{bmatrix}$$

### 1.1 Feed forward step

Consider the input  $x_1 = 3$ ,  $x_2 = 14$ . What is the final output  $(o_1, o_2)$  of the network?

```
[121]: import math
from math import exp
f = lambda x: max(0,x)
fexp = lambda input: map(lambda x: exp(f(x)), input)
softmax = lambda input: list(map(lambda x: x/sum(fexp(input)), fexp(input)))
inner_vv = lambda v1, v2: sum(map(lambda x, y: x*y, v1, v2))
inner_Mv = lambda M, v: [inner_vv(M_row, v) for M_row in M]
```

```

W = [[1, 0, -1], [0, 1, -1], [-1, 0, -1], [0, -1, -1]]
V = [[1,1,1,1,0], [-1,-1,-1,-1,2]]
x = [3, 14]

z = inner_Mv(W, [*x, 1])
fz = list(map(f, z))

u = inner_Mv(V, [*fz, 1])

print(softmax(u))

print("Alternatively:")
from functools import reduce
# print(f"o1 = e^{f(u[0])}/({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for
↪x in u))})")
# print(f"o2 = e^{f(u[1])}/({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for
↪x in u))})")

print(reduce(lambda x,y: x+'\n' + y, (f"o{the_index + 1} = e^{f(the_u)}/
↪({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for x in u))})" for
↪the_index, the_u in enumerate(u))))

```

```
[0.9999996940977731, 3.059022269256247e-07]
```

Alternatively:

```
o1 = e^15/(e^15+e^0)
```

```
o2 = e^0/(e^15+e^0)
```

## 1.2 Decision Boundaries

In this problem we visualize the “decision boundaries” in  $x$ -space, corresponding to the four hidden units. These are the lines in  $x$ -space where the values of  $z_1, z_2, z_3, z_4$  are exactly zero. Plot the decision boundaries of the four hidden units using the parameters of  $W$  provided above.

Enter below the area of the region of your plot that corresponds to a negative ( $< 0$ ) value for all of the four hidden units.

**Answer**

Line1:

$$x_1 W_{11} + x_2 W_{21} + W_{01} = 0$$

$$x_1 - 1 = 0$$

$$x_1 = 1$$

Line2:

$$\begin{aligned}
x_1 W_{12} + x_2 W_{22} + W_{02} &= 0 \\
x_2 - 1 &= 0 \\
x_2 &= 1
\end{aligned}$$

Line3:

$$\begin{aligned}
x_1 W_{13} + x_2 W_{23} + W_{03} &= 0 \\
-x_1 - 1 &= 0 \\
x_1 &= -1
\end{aligned}$$


Line4:

$$\begin{aligned}
x_1 W_{14} + x_2 W_{24} + W_{04} &= 0 \\
-x_2 - 1 &= 0 \\
x_2 &= -1
\end{aligned}$$

```
[191]: import matplotlib.pyplot as plt

plt.figure()
plt.plot([1, 1],[-1,1], 'b--', label='z1=0')
plt.arrow(1,0, 0.2* W[0][0], 0.2 * W[0][1], head_width=.04)
plt.plot([-1, 1],[1,1], 'r--', label='z2=0')
plt.arrow(0,1, 0.2* W[1][0], 0.2 * W[1][1], head_width=.04)
plt.plot([-1, -1],[-1,1], 'b-.', label='z3=0')
plt.arrow(-1,0, 0.2* W[2][0], 0.2 * W[2][1], head_width=.04)
plt.plot([-1, 1],[-1,-1], 'r-.', label='z4=0')
plt.arrow(0,-1, 0.2* W[3][0], 0.2 * W[3][1], head_width=.04)
plt.legend()
plt.show()
plt.close()

print('Area = 4 units')
```



Area = 4 units

### 1.3 Output of Neural Network

Using the same matrix  $V$  as above, what is the value of  $o_1$  (accurate to at least three decimal places if responding numerically) in the following three cases?

Assuming that  $f(z_1) + f(z_2) + f(z_3) + f(z_4) = 1$ :

```
[202]: sums = 1
u = [sums +0 , -sums +2]
print(reduce(lambda x,y: x+'\n' + y, (f"o{the_index +1} = e^{f(the_u)}/
↳({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for x in u))})" for_
↳the_index, the_u in enumerate(u))))
print('Values:', softmax(u))
```

$o_1 = e^1/(e^1+e^1)$

$o_2 = e^1/(e^1+e^1)$

Values: [0.5, 0.5]

Assuming that  $f(z_1) + f(z_2) + f(z_3) + f(z_4) = 0$ :

```
[203]: sums = 0
u = [sums +0 , -sums +2]
print(reduce(lambda x,y: x+'\n' + y, (f"o{the_index +1} = e^{f(the_u)}/
↳({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for x in u))})" for_
↳the_index, the_u in enumerate(u))))
```

```
print('Values:', softmax(u))
```

```
o1 = e^0/(e^0+e^2)
```

```
o2 = e^2/(e^0+e^2)
```

```
Values: [0.11920292202211755, 0.8807970779778824]
```

Assuming that  $f(z_1) + f(z_2) + f(z_3) + f(z_4) = 3$ :

```
[204]: sums = 3
u = [sums +0 , -sums +2]
print(reduce(lambda x,y: x+'\n' + y, (f"o{the_index +1} = e^{f(the_u)}/
↪({reduce(lambda x,y: x+'+' + y, ('e^' + str(f(x)) for x in u))})" for_
↪the_index, the_u in enumerate(u))))
print('Values:', softmax(u))
```

```
o1 = e^3/(e^3+e^0)
```

```
o2 = e^0/(e^3+e^0)
```


```
Values: [0.9525741268224333, 0.04742587317756678]
```

## 1.4 Inverse temperature

Just math.. fuck it

## 2 LSTM

The diagram below shows a single LSTM unit that consists of Input, Output, and Forget gates.



The behavior of such a unit as a recurrent neural network is specified by a set of update equations. These equations define how the gates, “memory cell”  $c_t$  and the “visible state”  $h_t$  are updated in response to input  $x_t$  and previous states  $c_{t-1}$ ,  $h_{t-1}$ . For the LSTM unit,

$$\begin{aligned}
 f_t &= \text{sigmoid}(W^{f,h}h_{t-1} + W^{f,x}x_t + b_f) \\
 i_t &= \text{sigmoid}(W^{i,h}h_{t-1} + W^{i,x}x_t + b_i) \\
 o_t &= \text{sigmoid}(W^{o,h}h_{t-1} + W^{o,x}x_t + b_o) \\
 c_t &= f_t \odot c_{t-1} + i_t \odot \tanh(W^{c,h}h_{t-1} + W^{c,x}x_t + b_c) \\
 h_t &= o_t \odot \tanh(c_t)
 \end{aligned}$$

where symbol  $\odot$  stands for element-wise multiplication. The adjustable parameters in this unit are matrices  $W^{f,h}$ ,  $W^{f,x}$ ,  $W^{i,h}$ ,  $W^{i,x}$ ,  $W^{o,h}$ ,  $W^{o,x}$ ,  $W^{c,h}$ ,  $W^{c,x}$  as well as the offset parameter vectors  $b_f$ ,  $b_i$ ,  $b_o$ , and  $b_c$ . By changing these parameters, we change how the unit evolves as a function of inputs  $x_t$ .

To keep things simple, in this problem we assume that  $x_t$ ,  $c_t$ , and  $h_t$  are all scalars. Concretely,

suppose that the parameters are given by:

$$\begin{array}{llll} W^{f,h} = 0 & W^{f,x} = 0 & b_f = -100 & W^{c,h} = -100 \\ W^{i,h} = 0 & W^{i,x} = 100 & b_i = 100 & W^{c,x} = 50 \\ W^{o,h} = 0 & W^{o,x} = 100 & b_o = 0 & b_c = 0 \end{array}$$

We run this unit with initial conditions  $h_{-1} = 0$  and  $c_{-1} = 0$ , and in response to the following input sequence:  $[0, 0, 1, 1, 1, 0]$  (For example,  $x_0 = 0$ ,  $x_1 = 0$ ,  $x_2 = 1$ , and so on).

## 2.1 LSTM states

Calculate the values  $h_t$  at each time-step and enter them below as an array  $[h_0, h_1, h_2, h_3, h_4, h_5]$ .

(Please round  $h_t$  to the closest integer in every time-step. If  $h_t = \pm 0.5$ , then round it to 0. For ease of calculation, assume that  $\text{sigmoid}(x) \approx 1$  and  $\tanh(x) \approx 1$  for  $x \geq 1$ , and  $\text{sigmoid}(x) \approx 0$  and  $\tanh(x) \approx -1$  for  $x \leq -1$ .)

```
[233]: Wfh = 0
Wih = 0
Woh = 0
Wfx = 0
Wix = 100
Wox = 100
b_f = -100
b_i = 100
b_o = 0
Wch = -100
Wcx = 50
b_c = 0

h_1 = 0
c_1 = 0
x = [0, 0, 1, 1, 1, 0]

sigmoid = lambda x: 1 if x>=1 else 0 if x<=-1 else 0.5*x + 0.5
tanh     = lambda x: 1 if x>=1 else -1 if x<=-1 else x

f = lambda ht_1, x_t, Wfx=Wfx, Wfh=Wfh, b_f=b_f: sigmoid(Wfh * ht_1 + Wfx *
    ↪ x_t + b_f)
i = lambda ht_1, x_t, Wix=Wix, Wih=Wih, b_i=b_i: sigmoid(Wih * ht_1 + Wix *
    ↪ x_t + b_i)
o = lambda ht_1, x_t, Wox=Wox, Woh=Woh, b_o=b_o: sigmoid(Woh * ht_1 + Wox *
    ↪ x_t + b_o)

c = lambda ft, it, ct_1, ht_1, x_t, Wch=Wch, Wcx=Wcx, b_c=b_c: ft * ct_1
    ↪ + it * tanh(Wch * ht_1 + Wcx * x_t + b_c)
h = lambda o_t, c_t: o_t * tanh(c_t)
```



```

ht_1 = [h_1]
ct_1 = [c_1]
for j in range(6):
    ft = f(ht_1[j], x[j])
    it = i(ht_1[j], x[j])
    ot = o(ht_1[j], x[j])
    ct_1.append( c(ft, it, ct_1[j], ht_1[j], x[j]) )
    ht_1.append( round(h(ot, ct_1[j+1])) )

print(ht_1[1:])

```

[0, 0, 1, -1, 1, 0]

## 2.2 LSTM states 2

Now, we run the same model again with the same parameters and same initial conditions as in the previous question. The only difference is that our input sequence is now: [1, 1, 0, 1, 1].

Calculate the values  $h_t$  at each time-step and enter them below as an array  $[h_0, h_1, h_2, h_3, h_4, h_5]$ .

```

[236]: x = [1,1,0,1,1]

ht_1 = [h_1]
ct_1 = [c_1]
for j in range(5):
    ft = f(ht_1[j], x[j])
    it = i(ht_1[j], x[j])
    ot = o(ht_1[j], x[j])
    ct_1.append( c(ft, it, ct_1[j], ht_1[j], x[j]) )
    ht_1.append( round(h(ot, ct_1[j+1])) )

print(ht_1[1:])


```

[1, -1, 0, 1, -1]

## 3 Backpropagation

One of the key steps for training multi-layer neural networks is stochastic gradient descent. We will use the back-propagation algorithm to compute the gradient of the loss function with respect to the model parameters.

Consider the  $L$ -layer neural network below:



In the following problems, we will use the following notation:  $b_j^l$  is the bias of the  $j^{\text{th}}$  neuron in the  $l^{\text{th}}$  layer,  $a_j^l$  is the activation of  $j^{\text{th}}$  neuron in the  $l^{\text{th}}$  layer, and  $w_{jk}^l$  is the weight for the connection from the  $k^{\text{th}}$  neuron in the  $(l-1)^{\text{th}}$  layer to the  $j^{\text{th}}$  neuron in the  $l^{\text{th}}$  layer.

If the activation function is  $f$  and the loss function we are minimizing is  $C$ , then the equations describing the network are:

$$a_j^l = f \left( \sum_k w_{jk}^l a_k^{l-1} + b_j^l \right)$$

$$\text{Loss} = C(a^L)$$

Note that notations without subscript denote the corresponding vector or matrix, so that  $a^l$  is activation vector of the  $l^{\text{th}}$  layer, and  $w^l$  is the weights matrix in  $l^{\text{th}}$  layer.

For  $l = 1, \dots, L$ .

### 3.1 Computing the Error

Let the weighted inputs to the  $d$  neurons in layer  $l$  be defined as  $z^l = w^l a^{l-1} + b^l$ , where  $z^l \in \mathbb{R}^d$ . As a result, we can also write the activation of layer  $l$  as  $a^l \equiv f(z^l)$ , and the “error” of neuron  $j$  in layer  $l$  as  $\delta_j^l \equiv \frac{\partial C}{\partial z_j^l}$ . Let  $\delta^l \in \mathbb{R}^d$  denote the full vector of errors associated with layer  $l$ .

Back-propagation will give us a way of computing  $\delta^l$  for every layer.

Assume there are  $d$  outputs from the last layer (i.e.  $a^L \in \mathbb{R}^d$ ). What is  $\delta_j^L$  for the last layer?

Loads of arithmetic later....